

COMP 737011 - Memory Safety and Programming Language Design

Lecture 2: Memory Allocation

徐 辉

xuh@fudan.edu.cn





Outline

- 1. Virtual Memory System
- 2. Dynamic Memory Allocation
- 3. Auto Memory Reclaim

1. Virtual Memory System

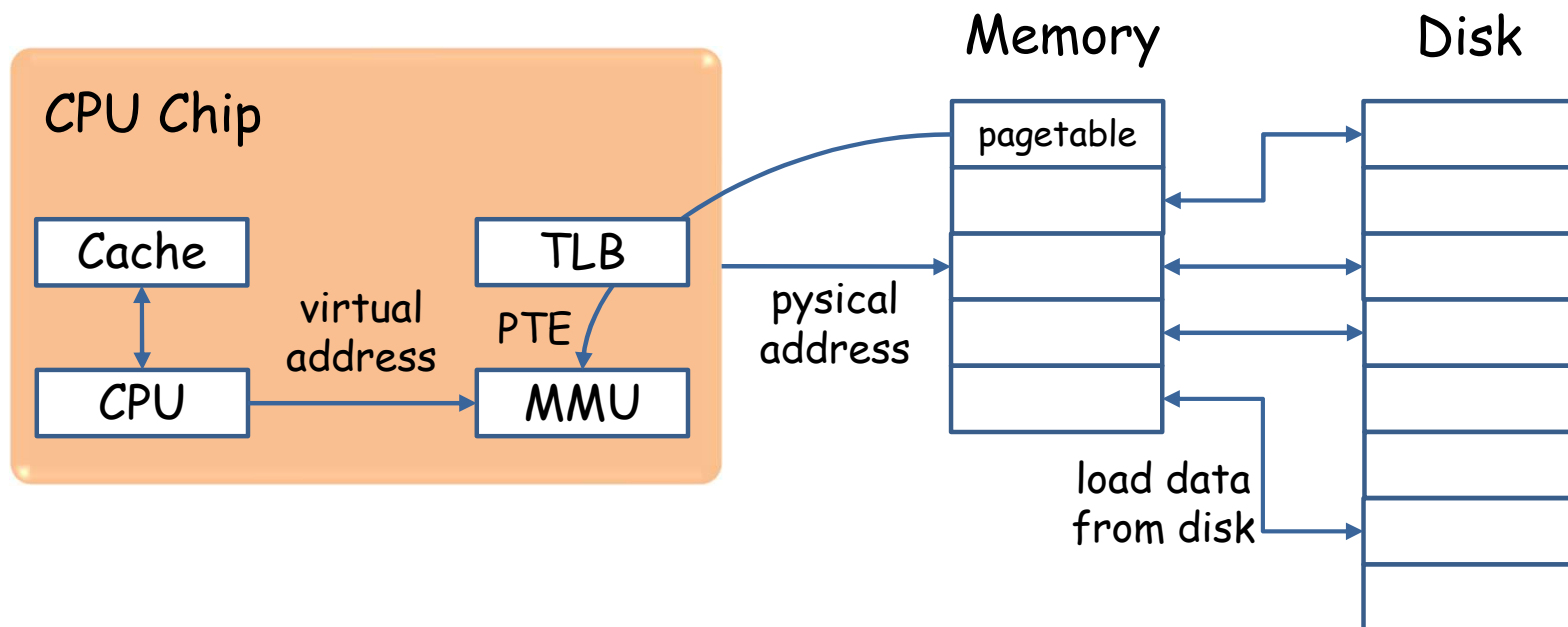
How to Support Multi-tasking OS?

- All processes share the same memory space?
 - Each process uses an exclusive memory region
 - e.g., unikernel or library OS
- Each process uses a distinct memory space
 - Virtual memory addressing
 - Both Linux and Windows use VM



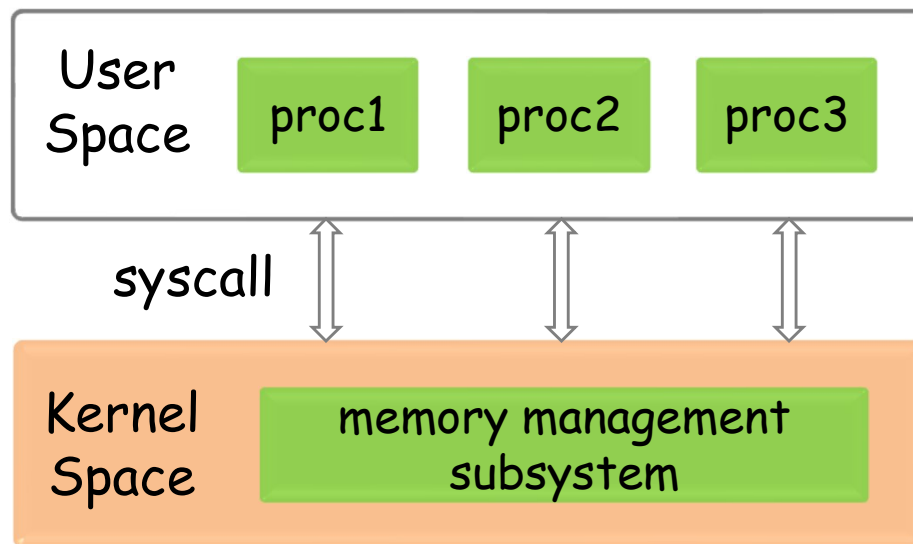
Virtual Memory Addressing

- MMU translates each virtual address to corresponding physical address by looking up the page table
- Cache page table entries with TLB
- Triggers page fault if the page is unavailable in DRAM



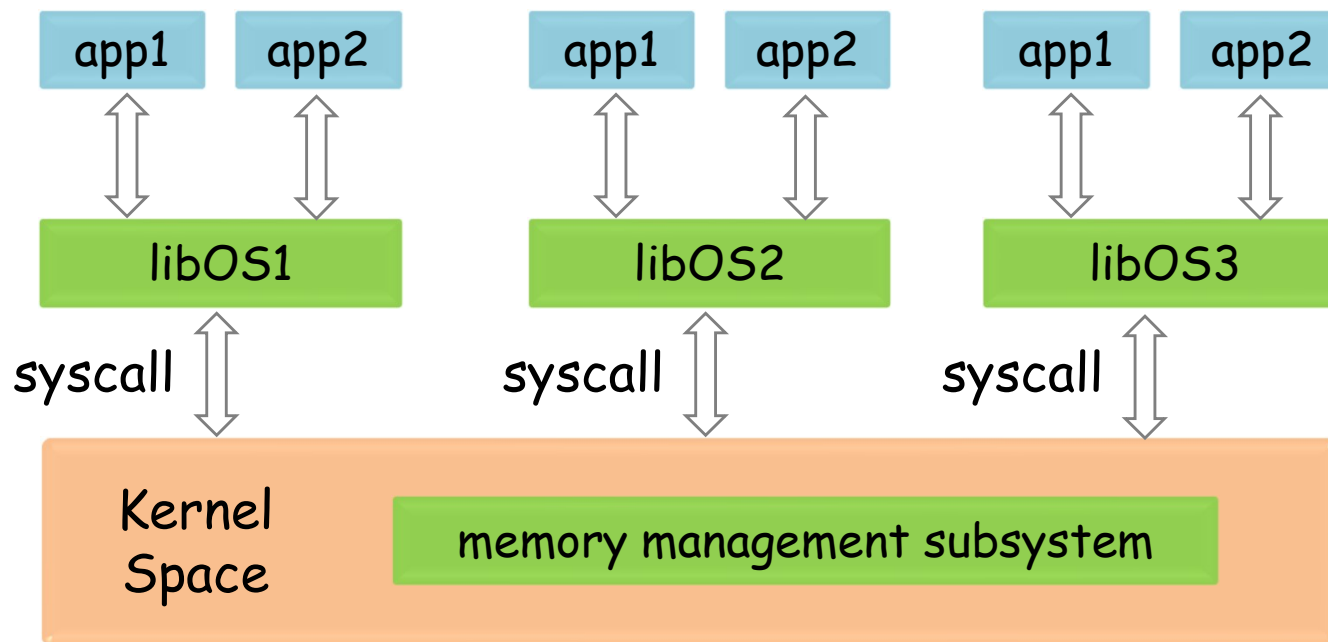
OS for VM

- Each process has a unique memory space
- Kernel responses for memory management
 - Map of memory space
 - Addressing
 - Handling page faults
- User space interacts with kernel via syscall



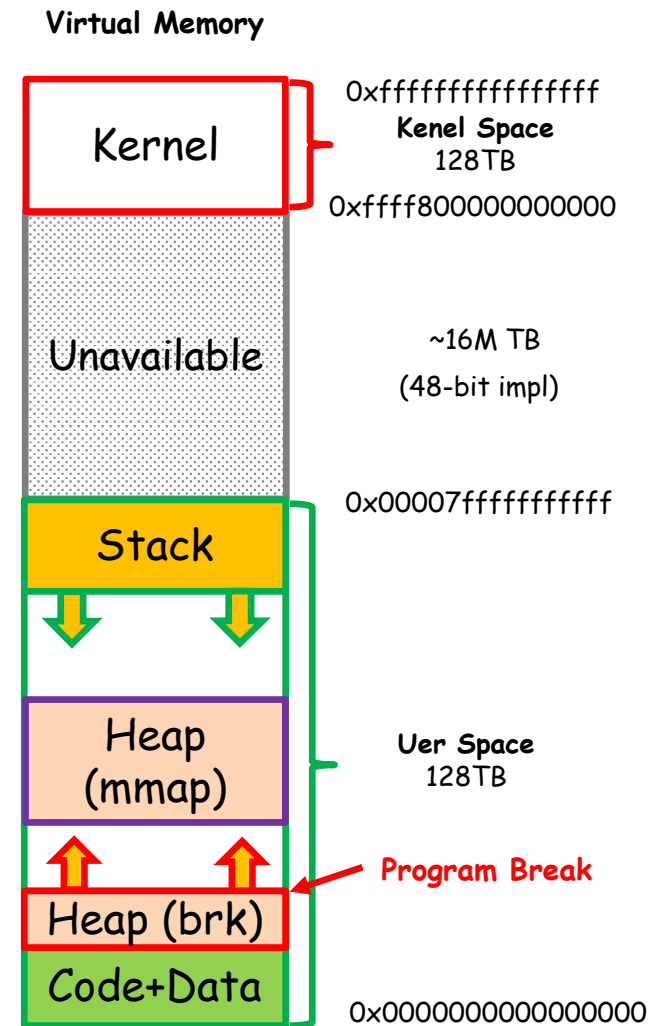
Unikernel

- Fault isolation for processes is difficult
 - Require instruction-level boundary checking
- Mainly used in cloud as libOS
 - Each user runs a libOS instead of a Docker image or VM



Memory Allocation in Linux

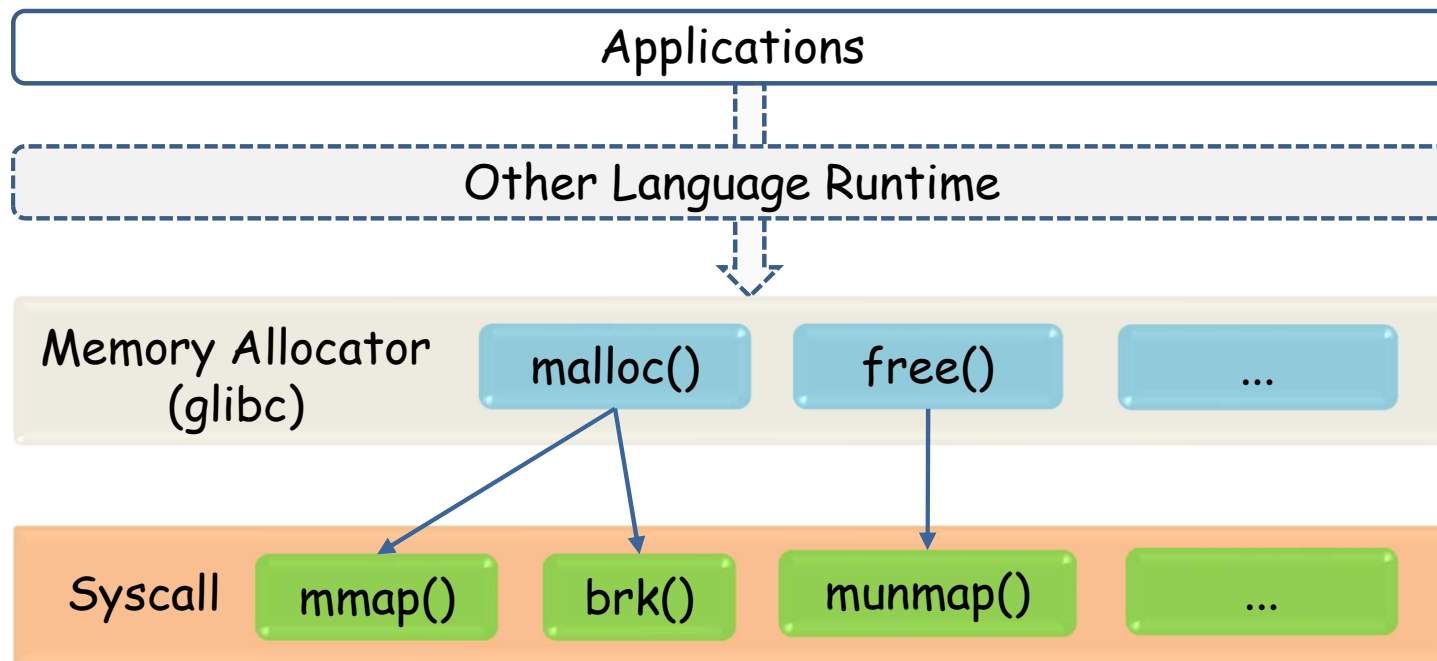
- Static allocation: static data
 - Compile-time constant
- Automatic allocation: stack
 - Each function has a stack frame
 - Multithreading program has multiple independent stacks
 - Compile-time constant
- Dynamic allocation: heap
 - More flexible



2. Dynamic Memory Allocation

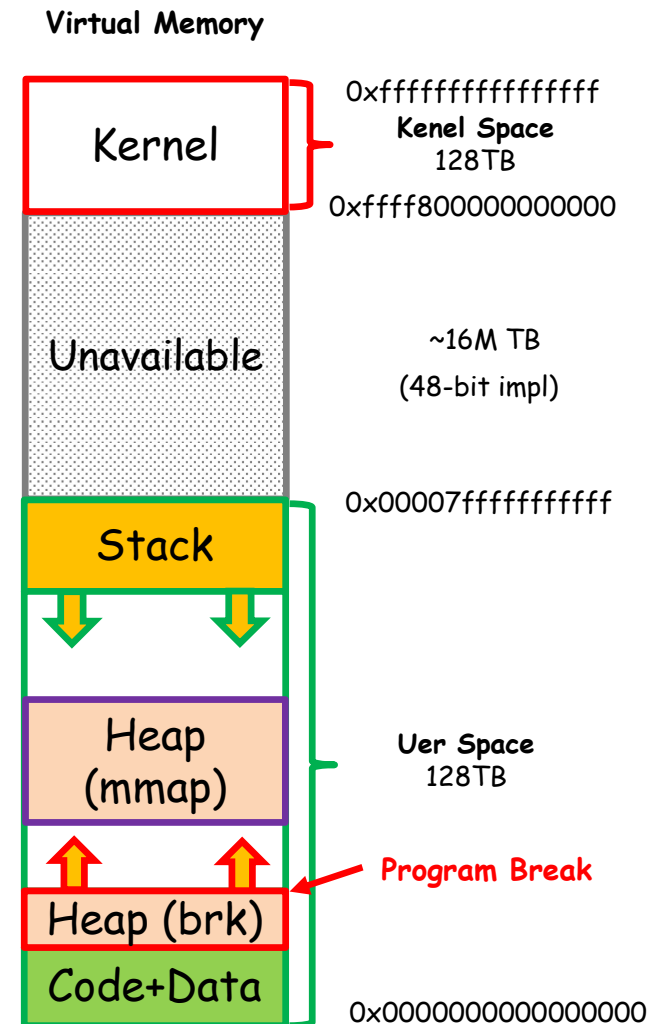
Overview of Memory Allocation APIs

- Applications use memory allocation APIs to acquire and release memory
- Memory allocator provides user-friendly APIs
 - e.g., `malloc()` and `free()` in glibc APIs
- Memory allocator invokes syscalls for achieving memory allocation.
 - e.g., `brk()` and `mmap()` in Linux



Heap Management in Linux

- Program break
 - Linux syscall brk()
 - For small-size memory trunks
 - Increase the brk pointer for memory allocation
 - Continuous address space
- Memory mapping:
 - Linux syscall mmap()
 - For file mapping and memory of large size (usually 256 KB)
 - Freed via munmap()



brk()/sbrk()/mmap()

```
int brk(void* end_data_segment); //Linux syscall
//change brk pointer to the specified addr value

void *sbrk(intptr_t increment); //a library API in Linux
//increase the brk pointer according to the increment

void *mmap(void *addr, // starting address
           size_t length, // byte
           int prot, // memory protection: read/write/exec
           int flags, // visibility: shared or private
           int fd, // file descriptor
           off_t offset); // offset of the file

int munmap(void *addr, size_t length);
```

glibc APIs for Heap Management

- `malloc(size_t n)`
 - Allocate a new memory space of size `n`
 - The memory is not cleared
 - Return the address pointer
- `free(void * p)`
 - Release the memory space pointed by `p`
 - Do not return to the system directly (for `brk`)
 - What would happen if `p` is null or already freed?
- `calloc(size_t nmemb, size_t size)`
 - Allocate an array of `nmemb * size` byte
 - The memory is set to zero
- `realloc(void *p, size_t size)`
 - Resize the memory block pointed by `p` to `size` bytes

Design Challenges for Allocator

- Each syscall costs nearly a hundred CPU cycles
 - =>An allocator should not frequently invoke syscalls
- Heap data are not compact
 - =>Reducing brk pointer for free is less effective
- How to manage and reuse freed memory chunks?

Basic Idea: Doug Lea's Allocator (dlmalloc)


- Freed memory chunks are managed as bins
- Each bin is a double-lined list of freed chunks of "fixed" size
- malloc() finds the corresponding bin for allocation
 - index is computed by logical shift: e.g., $\text{size} \gg 3$
 - first-in-first-out: each chunk with the same opportunity to be consolidated

Bins for sizes < 512 bytes contain chunks of all the same size, spaced 8 bytes apart. Larger bins are approximately logarithmically spaced:

large bins

small bins

64 bins of size	8
32 bins of size	64
16 bins of size	512
8 bins of size	4096
4 bins of size	32768
2 bins of size	262144
1 bin of size	what's left



The bins top out around 1MB because we expect to service large requests via mmap.

Structure of Chunks: Boundary Tag

- Sizes of free chunks are stored both in the front of each chunk and at the end.
- This makes consolidating fragmented chunks into bigger chunks very fast.

prev_size	
size	PREV_INUSE
forward pointer	
backward pointer	
unused space	
size	

free trunk

prev_size	
size	PREV_INUSE
user data	
size	

allocated trunk

Fastbins and Unsorted Bins


- Design consideration for consolidation is expensive
- Fastbins: light-weight bins in single-linked list
 - cannot be coalesced with adjacent chunks automatically
- Unsorted bins: free chunks are first put into unsorted bins

	list	coalesce	data
Fast bin	single-linked	no	small
Regular bin	double-linked	may	could be large

More Allocators

- ptmalloc (pthread malloc): used in glibc
 - a fork of dlmalloc with threading-related improvements
 - <https://sourceware.org/glibc/wiki/MallocInternals>
- tcmalloc (thread-caching malloc) by Google
 - <https://google.github.io/tcmalloc/>
- jemalloc
 - <http://jemalloc.net/>

3. Auto Memory Reclaim




Auto Reclaim Challenge

- Memory units allocated on stack are automatically reclaimed when a function returns
- Heap is hard to be reclaimed automatically
 - There could be multiple references across functions
 - Pointer analysis is NP-hard in general

Mechanisms for Auto Reclaim

- Mainly based on dynamic analysis
- Smart pointer
 - Dynamically track the number of references with a reference counter
 - Reclaim the memory once no variable owns it
- Garbage collection
 - Periodically check object references



Smart Pointer: unique_ptr

- Object is uniquely owned by one pointer
- User can transfer ownership through move()

```
int main() {  
    unique_ptr<MyClass> up1(new MyClass(2));  
    //unique_ptr<MyClass> up2 = up1; //compilation error  
    unique_ptr<MyClass> up2 = move(up1);  
    //cout << up1->val << endl; //segmentation fault  
    cout << up2->val << endl;  
}
```

Smart Pointer: `shared_ptr`

- Object are shared among pointers with a reference counter.
- Destroyed when the last remaining `shared_ptr` owning the object is destroyed or reassigned.

```
int main() {  
    shared_ptr<MyClass> sp1(new MyClass(2));  
    shared_ptr<MyClass> sp2 = p1;  
}
```

Guess the Output?


```
class MyClass{
public:
    int val;
    MyClass(int v) { val = v; }
    ~MyClass() { cout << "delete obj:"<< val << endl; }
};

int main() {
    MyClass* p0 = new MyClass(1);
    {
        shared_ptr<MyClass> p1(new MyClass(2));
        shared_ptr<MyClass> p2 = p1;
        shared_ptr<MyClass> p3(p0);
    }
    cout << p0->val << endl;
}
```

```
./a.out
delete obj:1
delete obj:2
0
```


Smart Pointer: weak_ptr

- Problem of shared_ptr
 - reference cycles
 - both p1 and p2 are leaked
- weak_ptr:
 - do up update the reference counter



```
class MyList{
public:
    int val;
    //shared_ptr<MyList> next;
    weak_ptr<MyList> next;
    ~MyList() { cout << "delete obj:"<< val << endl; }
};


int main() {
    shared_ptr<MyList> p1 = make_shared<MyList>();
    shared_ptr<MyList> p2 = make_shared<MyList>();
    p1->val = 1;
    p2->val = 2;
    p1->next = p2;
    p2->next = p1;
}
```

Garbage Collection

- When should the GC be triggered?
- What kind of objects should be recycled?
 - Reachability analysis
- How to recycle?
 - Slowdown due to intensive GC operation
 - Memory fragmentation issue


Reachability Analysis

- Stop the world
- Analyze from the root
- Unreachable objects should be recycled immediately



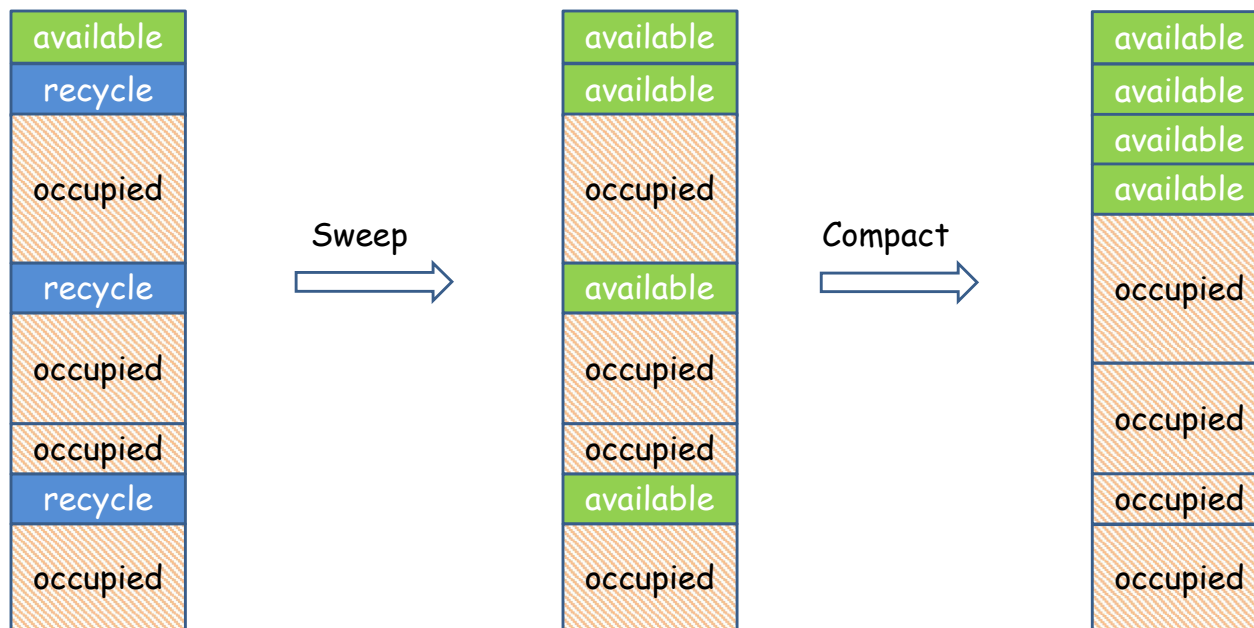
Incremental Analysis

- Do not need to stop the world
- Use three colors to record the temporary result
 - Orange: reached, and analysis (to other objects) is done
 - Gray: reached, but analysis is not finished
 - White: unreachable object
- false negative?
- false positive?



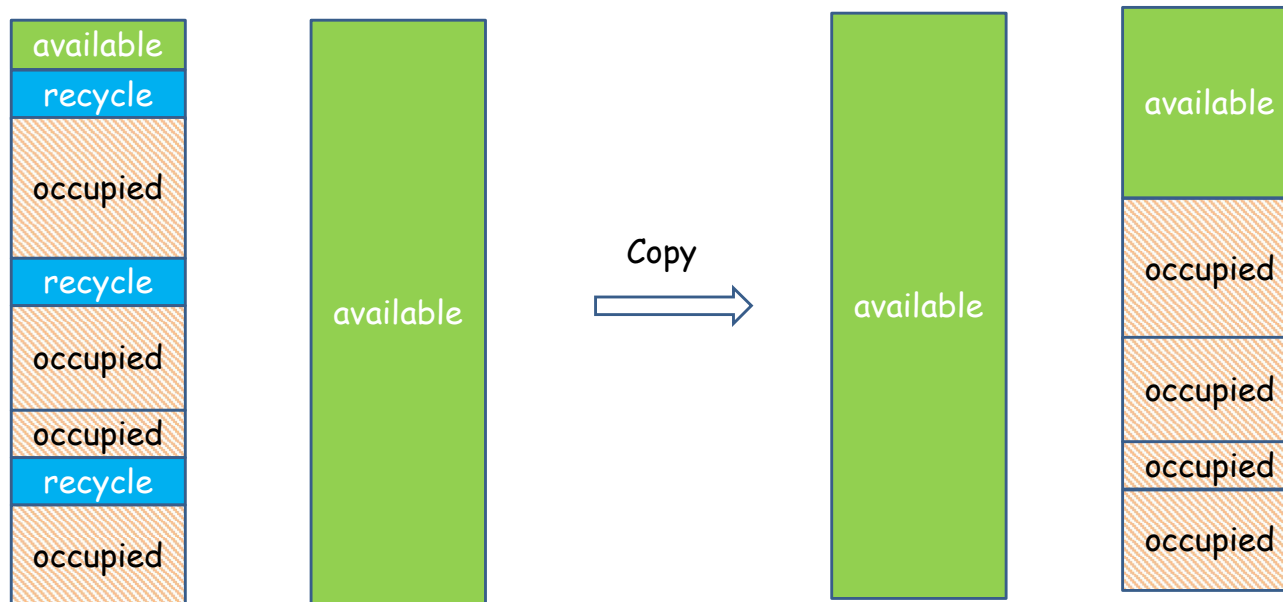
How to Recycle?

- For consecutive memory chunks (e.g., program break)
- Mark-sweep: suffers fragmentation issue
- Mark-compact: move all used units to one side
 - nontrivial overhead for moving data
 - when should the process be triggered?



Mark-Copy

- Two pieces of memory with the same size
 - the memory piece is still usable during copy
 - tradeoff between time and space




Observation

- Newly created objects tend to be recycled
- The objects survived after several GC rounds has a high chance to survive in the following round
- How can we utilize the observation for optimization?
 - Avoid frequent copy of old objects

Generational Collection

- Eden: for new objects
 - trigger minor GC if no space available
- Survivor: to host survived objects after minor GC
 - with two sub areas: from, to
 - Minor GC(eden+from) \Rightarrow to,
 - Minor GC(eden+to) \Rightarrow from
- Old: for objects survived after several rounds of minor GC
 - trigger major GC if no space available
 - large objects are saved to this area directly to avoid the overhead of copy.



Implementing GC for C?

- You may refer the following tutorials
 - <https://maplant.com/gc.html>
 - BoehmGC: <https://www.hboehm.info/gc/#details>

More Reference

- <https://sourceware.org/glibc/wiki/MallocInternals>
- https://cw.fel.cvut.cz/old/_media/courses/a4m33pal/04_dynamic_memory_v6.pdf
- https://heap-exploitation.dhavalkapil.com/diving_into_glibc_heap/bins_chunks