

COMP 737011 - Memory Safety and Programming Language Design

# Lecture 5: Concurrent Access

徐 辉

[xuh@fudan.edu.cn](mailto:xuh@fudan.edu.cn)



# Problem

- Risks of concurrent programs
  - data race or shared access
  - out-of-order execution
    - Compiler issue
    - CPU issue

# Outline

- 1. Data race and atomicity
- 2. Instruction reorder and memory barrier

# 1. Data Race and Atomicity

---

# Warm Up: Data Race

- Add operation is not atomic
  - load-add-store (multiple instructions or micro ops)

```
#define NUM 100
int global_cnt = 0;

void *mythread(void *in) {
    for (int i=0; i<NUM; i++)
        global_cnt++;
}

int main(int argc, char** argv) {
    pthread_t tid[NUM];
    for (int i=0; i<NUM; i++){
        assert(pthread_create(&tid[i], NULL, mythread, NULL)==0);
    }
    for (int i=0; i<NUM; i++){
        pthread_join(tid[i], NULL);
    }
    assert(global_cnt==NUM*NUM);
}
```

concurrently accessed by multiple threads

assertion could fail

Experimental hint: do not turn on optimization

# Atomic Version

way1: by declaring the variable as atomic

```
#define NUM 100
atomic_int global_cnt;

void *mythread(void *from) {
    // __atomic_fetch_add(&global_cnt, 1, __ATOMIC_SEQ_CST);
    for (int i=0; i<NUM; i++)
        global_cnt++;
}

int main(int argc, char** argv) {
    pthread_t tid[NUM];
    for (int i=0; i<NUM; i++){
        assert(pthread_create(&tid[i], NULL, mythread, NULL)==0);
    }
    for (int i=0; i<NUM; i++){
        pthread_join(tid[i], NULL);
    }
    assert(global_cnt==NUM*NUM);
}
```

way2: use atomic API

# Mechanism of Atomicity

- One instruction directly operates on the memory
  - do not load the variable to the register
- Lock prefix guarantees atomicity of Micro Ops
  - X86 provides a "lock" prefix for atomicity

```
gef> disass mythread
```

```
Dump of assembler code for function mythread:
```

```
0x00401150 <+0>:    push    rbp
0x00401151 <+1>:    mov     rbp, rsp
0x00401154 <+4>:    mov     QWORD PTR [rbp-0x10], rdi
0x00401158 <+8>:    mov     DWORD PTR [rbp-0x14], 0x0
0x0040115f <+15>:   cmp     DWORD PTR [rbp-0x14], 0x3e8
0x00401166 <+22>:   jge     0x401182 <mythread+50>
0x0040116c <+28>:   lock    add    DWORD PTR [rip+0x2ed0], 0x1
                                # <global_cnt>
0x00401174 <+36>:   mov     eax, DWORD PTR [rbp-0x14]
0x00401177 <+39>:   add     eax, 0x1
0x0040117a <+42>:   mov     DWORD PTR [rbp-0x14], eax
0x0040117d <+45>:   jmp     0x40115f <mythread+15>
0x00401182 <+50>:   mov     rax, QWORD PTR [rbp-0x8]
0x00401186 <+54>:   pop     rbp
0x00401187 <+55>:   ret
```

# Mutual Exclusion or Mutex

- How to achieve atomicity for a sequence of code?
  - entering the critical region without interference
- Approaches to implement a mutex
  - based on the cmpxchg instruction of intel CPU
  - based on Peterson algorithm



# CAS: Compare and Set/Swap

- How to achieve Atomic CAS?
  - x86 instruction: `cmpxchg`
  - C API: `atomic_compare_exchange_strong`

#based on rax

`lock cmpxchg dst src`

## Semantic:

```
if(dst == eax) { dst = src; ZERO_FLAG = 1; }  
else { eax = dst; ZERO_FLAG = 0; }
```

`atomic_compare_exchange_strong(&dst, &test, src)`

exactly the same with `cmpxchg`

# Peterson Algorithm for Mutex

```
void* t0(void *from) {  
    flag[0] = true;  
    turn = 1;  
    while(flag[1]==true && turn==1)  
        sleep(1);  
    do_critical();  
    flag[0] = false;  
}
```

```
void* t1(void *from) {  
    flag[1] = true;  
    turn = 0;  
    while(flag[0]==true && turn==0)  
        sleep(1);  
    do_critical();  
    flag[1] = false;  
}
```

```
int flag[2], turn;  
int x=0;  
  
void do_critical(){  
    x++;  
}  
  
int main(int argc, char** argv) {  
    pthread_t tid[2];  
    assert(pthread_create(&tid[0], NULL, t0, NULL)==0);  
    assert(pthread_create(&tid[1], NULL, t1, NULL)==0);  
    pthread_join(tid[0], NULL);  
    pthread_join(tid[1], NULL);  
}
```

# Question

- Is `shared_ptr` of C++ thread-safe?
  - reference counter
  - data read/write
- Can you implement a thread-safe construct?
- We will learn `ARC<Mutex<T>>` in Rust

## 2. Instruction Reorder & Memory Barrier

---

# Out-of-Order Execution

- Compiler reordering during optimization
- CPU out-of-order execution
- This lecture focuses on compile-time ordering

# Compiler Reordering

- Suppose optimization (e.g., -O2 or O3) is enabled, the compiler might make mistakes.

```
atomic_int a = 1;
```

```
void *t0 (void* in){  
    while (a);  
}
```

```
void *t1 (void* in){  
    a = 0;  
}
```

```
0x00401150 <+0>:    cmp     DWORD PTR [rip+0x2ee9],0x0    # <a>  
0x00401157 <+7>:    je      0x401162 <t0+18>  
0x00401159 <+9>:    nop     DWORD PTR [rax+0x0]  
0x00401160 <+16>: jmp     0x401160 <t0+16>  
0x00401162 <+18>: ret
```



infinite loop

# Another Example

- The following assertion could fail on some platforms if the execution order cannot be guaranteed

```
atomic_int a = 1;  
atomic_int b = 1;
```

```
void *t0 (void* in){  
    a = 0;  
    b = 0;  
}
```

```
void *t1 (void* in){  
    while(!b);  
    assert(!a);  
}
```

Note: the assertion does not fail on X86 but may fail on other platforms. Please refer to <https://stackoverflow.com/questions/48139399/memory-order-relaxed-not-work-as-expected-in-code-from-c-concurrency-in-action>

# Use Memory Barrier (Fence)

- Discard all variable values on registers
  - Reload them from memory
- Guarantee happens-before: operations prior to the barrier are always executed before operations after the barrier.

```
#define barrier() __asm__ __volatile__("" : : : "memory");
```

```
void *t0 (void* in){  
    while (a)  
        barrier();  
}
```

```
void *t0 (void* in){  
    a = 0;  
    barrier();  
    b = 0;  
}
```



# Relax the Variables

- We only want some variables to be updated:
  - use volatile when declaring a variable
- We only want the variable to be updated in specific program points:
  - Use ACCESS\_ONCE(), which is also based on volatile
- Further relax the restrictions based on operations:
  - READ\_ONCE and WRITE\_ONCE()

```
volatile int a = 1;
void *t0 (void* in){
    while (a) ;
}
```

```
#define ACCESS_ONCE(x) (*(volatile typeof(x) *)&(x))
volatile int a = 1;
void *t0 (void* in){
    while (ACCESS_ONCE(a)) ;
}
```

# Relax Happens-Before Requirement

- Use specific memory ordering
  - Sequential consistency (default on x86)
    - the most strong one, no reordering across the barrier;
  - Acquire-release
    - release: no reads or writes in the current thread can be reordered after this store
    - acquire: no reads or writes in the current thread can be reordered before this load
    - commonly used for locks
  - Relaxed
    - no synchronization or ordering constraints
    - only atomicity

# Summary



# In-Class Practice

- Implement a mutex (lock) in C and demonstrate the effectiveness
  - based on the `atomic_compare_exchange_weak()` API, [https://en.cppreference.com/w/c/atomic/atomic\\_compare\\_exchange](https://en.cppreference.com/w/c/atomic/atomic_compare_exchange)
- Optional: try to implement a thread-safe shared pointer for some structs with mutex member functions.

# More Reference

- [https://www.alibabacloud.com/blog/memory-model-and-synchronization-primitive---part-1-memory-barrier\\_597460](https://www.alibabacloud.com/blog/memory-model-and-synchronization-primitive---part-1-memory-barrier_597460)
- <https://gcc.gnu.org/wiki/Atomic/GCCMM/AtomicSync>
- [https://en.cppreference.com/w/cpp/atomic/memory\\_order](https://en.cppreference.com/w/cpp/atomic/memory_order)
- [https://www4.cs.fau.de/Lehre/WS17/V\\_CS/Uebungen/aufgaben/slides3.pdf](https://www4.cs.fau.de/Lehre/WS17/V_CS/Uebungen/aufgaben/slides3.pdf)
- [https://www.cs.cmu.edu/~410-s05/lectures/L31\\_LockFree.pdf](https://www.cs.cmu.edu/~410-s05/lectures/L31_LockFree.pdf)