

COMP 737011 - Memory Safety and Programming Language Design

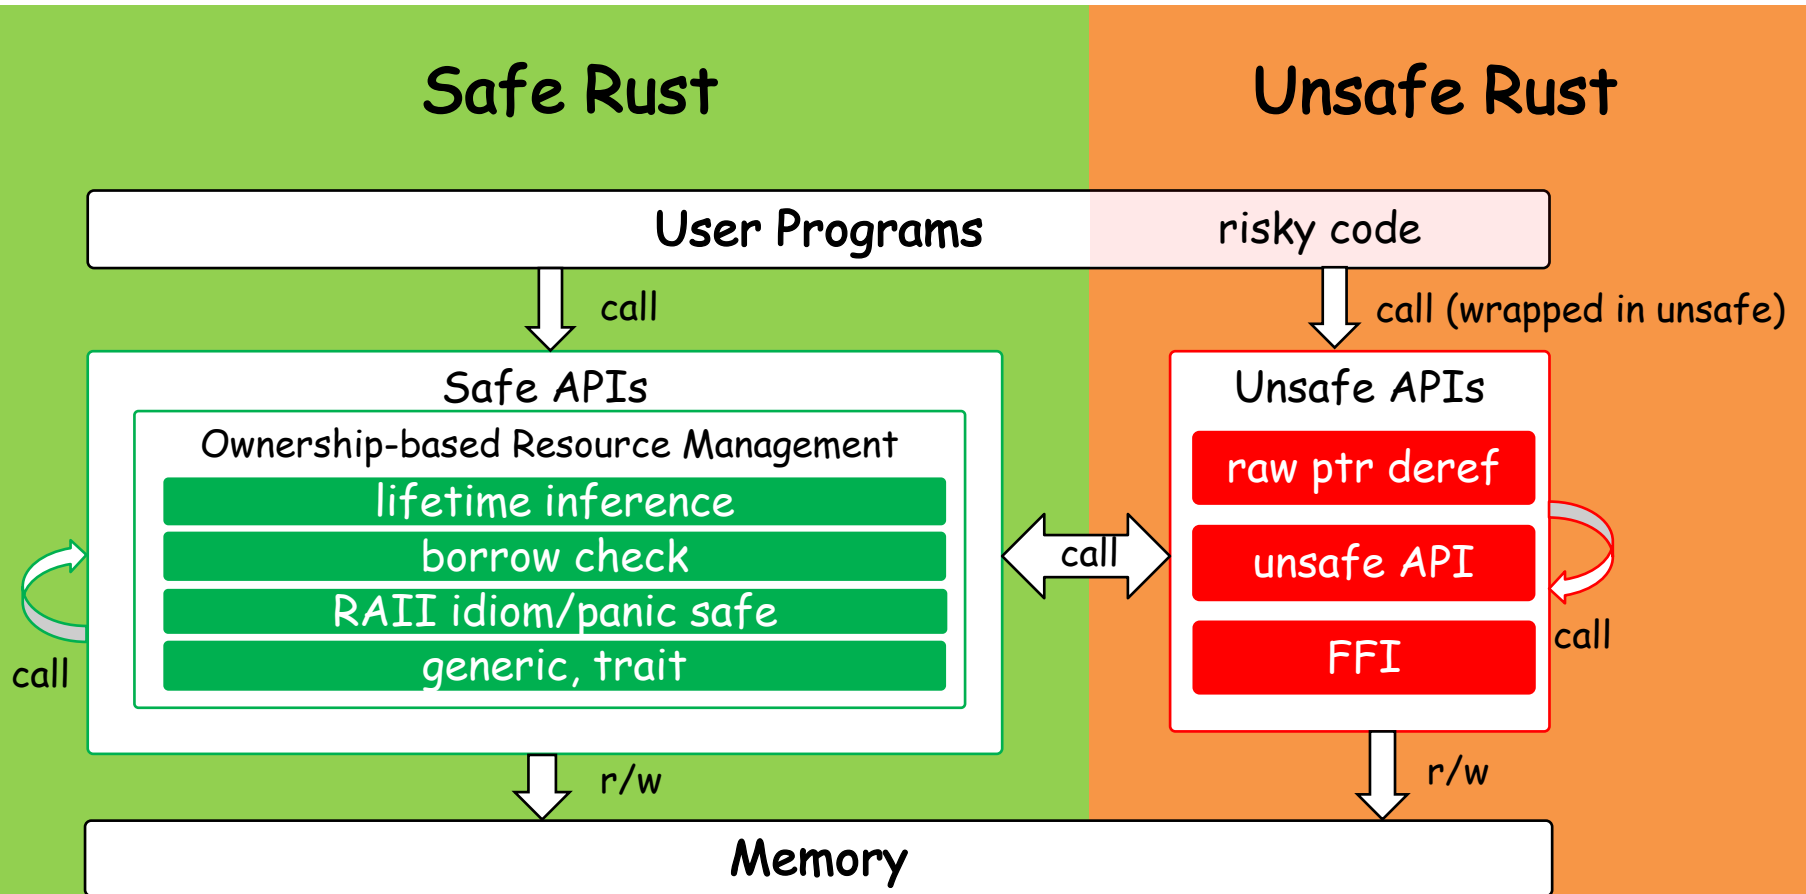
Lecture 6: Rust OBRM

徐 辉

xuh@fudan.edu.cn



Rust Overview



Outline

- 1. Ownership
- 2. RAII and Lifetime
- 3. Unsafe Code

1. Ownership

Motivation of Design

- Dangling pointer is unacceptable
- Causal of dangling pointer?
 - memory reclaim or object destruction
 - manual reclaim of heap data
 - we should prevent such manual reclaim
 - automatic reclaim => garbage collection or else?
 - stack object destruction
 - there could be other aliases to part of the object
 - alias analysis is NP-hard
 - shared pointer? inefficient
- Rust comes to rescue

Idea

- Ownership: each object is owned by one variable.
- Borrow: the ownership can be borrowed by other variables in two mode:
 - immutable: read only
 - mutable: read/write
- Key rules: exclusive mutability
 - two variables cannot share mutable access to the object at the same time.
- Benefit: make the alias analysis problem much easier, why?
 - only mutable pointers can lead to dangling pointers
 - we only need to trace the mutable pointer for each object

Ownership & Borrowing

- Borrowed ownership will be returned automatically if no longer used.

```
fn main(){  
    let mut alice = Box::new(1);  
    let bob = alice;  
    println!("bob:{}", bob);  
    println!("alice:{}", alice);  
}
```



alice owns the object

transfer the vector to bob
alice loses the ownership

```
fn main(){  
    let mut alice = Box::new(1);  
    let bob = &alice;  
    println!("bob:{}", bob);  
    println!("alice:{}", alice);  
}
```



bob borrows the ownership

bob returns the ownership
to alice automatically

Move Operator (=)

- If a type is not Copy (trait), move transfers the ownership.
 - e.g., Box<T> is not copy
- If a type is Copy, move does not transfer the ownership and only copies the value

```
fn main(){  
    let mut alice = 1;  
    let bob = alice;  
    println!("bob:{}", bob);  
    println!("alice:{}", alice);  
}
```

→ alice owns the object

→ copy the object to bob



Which Type Can be Copy?

- Primitive types on stack
- Composite types with all fields implementing copy
- How to (deep) copy objects of other non-Copy types:
 - implement Clone (trait)
 - each Copy type is also Clone.

```
fn main(){  
  let mut alice = Box::new(1);  
  let bob = alice.clone();  
  println!("bob:{}", bob);  
  println!("alice:{}", alice);  
}
```

→ alice owns the object


→ clone the object for bob




Mutability

alice is mutable

```
let mut alice = 1;  
alice+=1;
```




```
let alice = 1;  
alice+=1;
```




mutable borrow


```
let mut alice = 1;  
let bob = &mut alice;  
*bob+=1;
```



```
let mut alice = 1;  
let bob = &alice;  
*bob+=1;
```




```
let alice = 1;  
let bob = &mut alice;  
*bob+=1;
```




bob is mutable

```
let mut alice = 1;  
let mut carol = 1;  
let mut bob = &mut alice;  
*bob+=1;  
bob = &mut carol;  
*bob+=1;
```



```
let mut alice = 1;  
let mut carol = 1;  
let bob = &mut alice;  
*bob+=1;  
bob = &mut carol;  
*bob+=1;
```



What if...

```
fn main(){  
    let mut alice = 1;  
    let bob = &mut alice;  
    println!("bob:{}", bob);  
    println!("alice:{}", alice);  
}
```



mutable borrow

bob returns the ownership

```
fn main(){  
    let mut alice = 1;  
    let bob = &mut alice;  
    println!("alice:{}", alice);  
    println!("bob:{}", bob);  
}
```



exclusive mutability

Pros and Cons

- Benefit
 - Compile-time prevention of shared mutable aliases.
- When Shared Mutability is a Must...
 - Such as double linked list?
 - Two options (we will discuss later):
 - use shared pointer (reference counter)
 - unsafe code

2. RAII and Lifetime

Resource Acquisition is Initialization

Idea of RAII

- Ties resources to object lifetime
- Resource allocation is done during object creation by the constructor
 - all pointers refer to specific objects
 - no raw or dangling pointers
 - even no uninitialized memory
- Resource deallocation is done during object destruction by the destructor
 - no manual deallocation is needed
 - achieved through static lifetime inference

Lifetime

- Each object has a lifetime constraint
- The object is reclaimed automatically after death
- A variable cannot borrow an object with a shorter lifetime

```
fn main(){  
    let alice;  
    {  
        let bob = 5;  
        alice = &bob;  
    }  
    println!("alice:{}", alice);  
}
```



Review Move for Non-Copy Types

- Extend the lifespan of the object on heap

```
fn test(){  
    let alice;  
    {  
        let bob = Box::new(1);  
        alice = bob;  
    }  
    println!("alice:{}", alice);  
}
```

```
fn testret() -> Box<u64>{  
    Box::new(1)  
}
```

```
let r = testret();  
println!("return:{}", r);
```


Lifetime Declaration

- Lifetime constraint can be lexical during function declaration.

```
fn stringcmp(){
    let str1 = String::from("alice");
    let str2 = String::from("bob111");
    let result = longer(&str1, &str2);
    println!("The longer string is {}", result);
}

fn longer<'a>(x:&'a String, y:&'a String)->&'a String{
    if x.len()>y.len(){
        x
    } else {
        y
    }
}
```

Partial Order of Lifetime

- `<'a: 'b, 'b>` means lifetime `a > b`

```
fn stringcmp(){  
    let str1 = String::from("alice");  
    let result;  
    //{  
        let str2 = String::from("bob111");  
        result = longer(&str1, &str2);  
    //}  
    println!("The longer string is {}", result);  
}
```

```
fn longer<'a: 'b, 'b>(x:&'a String, y:&'b String)->&'b String{  
    if x.len()>y.len(){  
        x  
    } else {  
        y  
    }  
}
```

Non-lexical Lifetime

- The default mode is non-lexical unless necessary
- Rust compiler tries to minimize the lifespan

```
'a: { let str1 = "alice";  
    'b: { let str2 = "bob";  
        'c: { let result = longer(str1, str2);  
              println!("The longer string is {}", result);  
        }  
    }  
}
```

Lifetime Elision to Be More Ergonomic

- Sometimes lifetime declaration can be elided
 - exactly one input lifetime position
 - multiple positions, but one is `&self` or `&mut self`
- Assign the lifetime to elided output lifetimes.

```
substr<'a>(s: &'a str, until: usize) -> &'a str;  
fn substr(s: &str, until: usize) -> &str; // elided fn  
  
fn frob(s: &str, t: &str) -> &str; // ILLEGAL
```

More About Lifetime

- Static, e.g., all `str`
- Unbounded

```
let s: &str = "hello world";
```

is equivalent to the following one

```
let s: &'static str = "hello world";
```

```
fn get_str<'a>() -> &'a str;
```

Constructor

- Use struct/enum to define a new type
- Create an instance of a user-defined type by explicitly passing values to each field

```
struct MyType1 {  
    a:u8  
    b:u64,  
}  
let v = MyType1 {a:1, b:2};
```

```
enum MyType3{  
    a(u8),  
    b(u64),  
}  
let v = MyType1::a(1);
```

```
struct MyType2 {  
    a:u8  
    b:Box<u64>,  
}  
let v = MyType2 {a:1, Box::new(2)};
```

Automatic Reclaim

- Objects of Copy type (on stack) can be reclaimed automatically
- For other objects with heap data
 - Drop (trait) unused objects by calling the destructor
- Drop and Copy are exclusive in Rust
- Box<T> is Drop trait

```
struct Droppable{}

impl Drop for Droppable {
    fn drop(&mut self){
        println!("dropping...");
    }
}

fn testdrop(){
    let mut alice = Droppable{};
}
```

Recursive Drop

- Recursively call the destructor of each field
- Rust prevent manually calling `Drop::drop()`
- `mem::drop()` is another different function
 - can invoke `Drop::drop()` by consumes the ownership

```
struct Droppable{ }
struct ParentDrop{ a:Droppable, b:Droppable,}

impl Drop for Droppable {
    fn drop(&mut self){
        println!("dropping...");
    }
}

impl Drop for ParentDrop {
    fn drop(&mut self){
        println!("parent dropping...");
    }
}

let mut alice = ParentDrop{a:Droppable{},b:Droppable{}};
```


Some Limitations of So Far...

- RAII prevents uninitialized data, but sometimes uninitialized objects are needed
 - e.g., linked list
 - we can append the list dynamically
 - how to represent the next pointer?
- Shared mutable aliases are needed
 - e.g., double-linked list

Use Options for Uninitialized Objects?

- Options: an enumeration type
 - Some(T): the object type
 - None: if the object is uninitialized

```
pub enum Option<T> {  
    None,  
    Some(T),  
}  
  
let v = Some(...)  
match v.next {  
    Some(n) => ...,  
    None => panic!(),  
}
```

Example with a Singly-linked List

```
struct List{
    val: u64,
    next: Option<Box<List>>,
}
fn standard(){
    let mut l = List{val:1, next:None};
    l.next = Some(Box::new(List{val:2, next:None}));
    match l.next {
        Some(ref mut n) =>
            n.next = Some(Box::new(List{val:3, next:None})),
        None => panic!(),
    }
    let mut h = &l;
    loop {
        println!("{}", h.val);
        match h.next {
            Some(ref n) => h = n,
            None => break,
        }
    }
}
```

RAII for Thread Panic

- In a multi-threaded application, what happens when one thread exit exceptionally ?
 - abort: directly terminate the thread
 - panic: perform stack unwinding before exit
- Importance of RAII during stack unwinding
 - release locks (mutex)
 - release opened file descriptors
 - release allocated memories on heap

Sample Multi-threaded Program

- When a spawned thread panics, unwind its stack
- The main thread continues execution
- Ineffective for fatal errors: e.g., stack overflow

```
fn main() {  
    let handle = thread::spawn(|| {  
        for i in 0..5 {  
            println!("new thread print {}", i);  
            thread::sleep(Duration::from_millis(10));  
        }  
        panic!();  
        //recursive();  
    });  
    for i in 0..10 {  
        println!("main thread print {}", i);  
        thread::sleep(Duration::from_millis(10));  
    }  
    handle.join();  
}
```

3. Unsafe Rust

Unsafe

- Dereference raw pointers
- Call unsafe functions
- Call functions of foreign language (FFI)

```
let mut num = 5;
let r1 = &num as *const i32;
unsafe {
    println!("r1 is: {}", *r1);
}
```

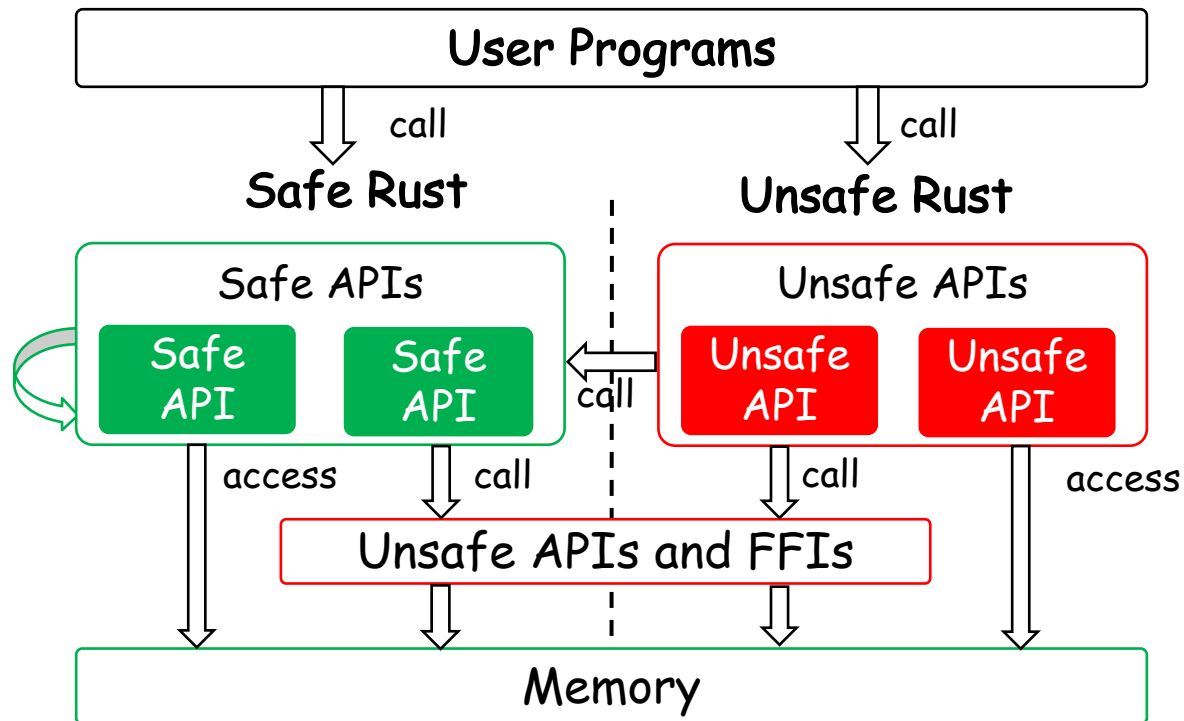
Dereference raw pointers

```
unsafe fn risky() {
    let address = 0x012345usize;
    let r = address as *const i32;
}
unsafe {
    risky();
}
```

Call unsafe functions


Principle of Rust

- Safe API should not incur undefined behaviours
- Interior unsafe: wrap unsafe code into safe APIs
- Avoid using unsafe code unless necessary.



Trust Model

- Rust does not trust developers, so only safe code is allowed
- If the developer declares he knows the risk, Rust will trust him



Problem for Double-Linked List

- Both prev and next owns a node within the list
- Violate exclusive mutability



```
struct List{  
    val: u64,  
    prev: Option<Box<List>>,  
    next: Option<Box<List>>,  
}
```

Raw Pointer Version (Not Recommend)

- The resource may not be dropped automatically
- Prone to dangling pointers

```
struct List{
    val: u64,
    next: *mut List,
    prev: *mut List,
}

fn rawptr(){
    let mut l = List{val:1, next:null_mut(), prev:null_mut()};
    l.next = &mut List{val:2, next:null_mut(), prev:null_mut()};
    unsafe {
        let mut cur = &mut *(l.next);
        cur.prev = &mut l;
        cur.next = &mut List{val:3, next:null_mut(), prev:null_mut()};
        (*(cur.next)).prev = cur;
    }
}
```

Solution Hint with RC and RefCell

- RC: single-threaded reference-counting pointer
 - RC enables shared immutable aliases
- RefCell: a mutable memory location with dynamically checked borrow rules

```
struct List{  
    val: u64,  
    prev: Option<Rc<RefCell<List>>>,  
    next: Option<Rc<RefCell<List>>>,  
}
```

RC

- Reference counter for shared aliases
- Mutate via `get_mut()`
 - mutual exclusion during compile time
 - if cloned, `get_mut()` returns `None` during run time

```
fn main(){  
    let mut x = Rc::new(1);  
    //let _y = Rc::clone(&x);  
    let t1 = Rc::get_mut(&mut x).unwrap();  
    //let t2 = Rc::get_mut(&mut x).unwrap();  
    *t1 = 2;  
    assert_eq!(*x, 2);  
  
    let _y = Rc::clone(&x);  
    assert!(Rc::get_mut(&mut x).is_none());  
}
```

if cloned, `get_mut()`
returns `None`

compile error




RefCell

- Perform borrow check during runtime


```
fn testrefcell(){  
    let x = RefCell::new(Box::new(1));  
    {  
        let mut y = x.borrow_mut();  
        //let z = x.borrow_mut();  
        *(*y) = 2;  
    }  
    assert_eq!(2, *(*x.borrow()));  
}
```

→ panic during runtime

RC vs RefCell



get_mut() may return None
if cloned



the second borrow_mut()
triggers panic

In-class Practice

- Implement a binary search tree with
 - insert function
 - search function
- Implement a double linked list with Safe Rust

Reference

- <https://doc.rust-lang.org/book/>
- <https://doc.rust-lang.org/stable/nomicon/>