

COMP 737011 - Memory Safety and Programming Language Design

Lecture 8: Rust Concurrency

徐辉

xuh@fudan.edu.cn



Module `std::sync`

Modules

`atomic` Atomic types
`mpsc` Multi-producer, single-consumer FIFO queue communication primitives.

Structs

`Arc` A thread-safe reference-counting pointer. ‘Arc’ stands for ‘Atomically Reference Counted’.

`Barrier` A barrier enables multiple threads to synchronize the beginning of some computation.

`BarrierWaitResult` A `BarrierWaitResult` is returned by `Barrier::wait()` when all threads in the `Barrier` have rendezvoused.

`Condvar` A Condition Variable

`Mutex` A mutual exclusion primitive useful for protecting shared data

`MutexGuard` An RAII implementation of a “scoped lock” of a mutex. When this structure is dropped (falls out of scope), the lock will be unlocked.

`Once` A synchronization primitive which can be used to run a one-time global initialization. Useful for one-time initialization for FFI or related functionality. This type can only be constructed with `Once::new()`.

`OnceState` State yielded to `Once::call_once_force()`’s closure parameter. The state can be used to query the poison status of the `Once`.

`PoisonError` A type of error which can be returned whenever a lock is acquired.

`RwLock` A reader-writer lock

`RwLockReadGuard` RAII structure used to release the shared read access of a lock when dropped.

`RwLockWriteGuard` RAII structure used to release the exclusive write access of a lock when dropped.

`WaitTimeoutResult` A type indicating whether a timed wait on a condition variable returned due to a time out or not.

`Weak` `Weak` is a version of `Arc` that holds a non-owning reference to the managed allocation. The allocation is accessed by calling `upgrade` on the `Weak` pointer, which returns an `Option<Arc<T>>`.

<https://doc.rust-lang.org/std/sync/index.html#structs>

Module `std::marker`

Structs

- PhantomData** Zero-sized type used to mark things that “act like” they own a `T`.
- PhantomPinned** A marker type which does not implement `Unpin`.

Traits

- DiscriminantKind** Experimental Compiler-internal trait used to indicate the type of enum discriminants.
- StructuralEq** Experimental Required trait for constants used in pattern matches.
- StructuralPartialEq** Experimental Required trait for constants used in pattern matches.
- Unsize** Experimental Types that can be “unsized” to a dynamically-sized type.
- Copy** Types whose values can be duplicated simply by copying bits.
- Send** Types that can be transferred across thread boundaries.
- Sized** Types with a constant size known at compile time.
- Sync** Types for which it is safe to share references between threads.
- Unpin** Types that can be safely moved after being pinned.

Outline

- 1. Basic features
 - a) Atomic types
 - b) ARC/Weak
 - c) Memory barrier
 - d) Locks: Mutex/RwLock
 - e) Conditional variable
 - f) Once
 - g) mpvc
- 2. Marker Trait
 - Send and Sync

1. Basic features

a) Atomic Types

- Several atomic types
 - AtomicBool,
 - AtomicIsize,
 - AtomicUsize,
 - ...
- Similar to C++ `std::atomic`

```
let mut foo = AtomicI32::new(0);  
*foo.get_mut() = 5  
foo.fetch_add(10, Ordering::SeqCst);  
foo.compare_and_swap(5, 10, Ordering::Relaxed);
```

_____ assignment
_____ atomic add
_____ CAS

b) Arc<T>: Atomically Ref Counted

- Similar to RC<T>, but is thread safe
- Use atomic operations for reference counting
- Mutating through an Arc generally use Mutex, RwLock, etc.

```
fn main() {  
    let v = Arc::new(Mutex::new(vec![1,2,3]));  
    for i in 0..3 {  
        let cloned_v = v.clone();  
        thread::spawn(move || {  
            cloned_v.lock().unwrap().push(i);  
        });  
    }  
}
```

move the ownership to
the thread closure



c) Memory Barrier

- Ordering is an enumerate type
 - SeqCst: sequential consistency
 - Acquire-Release: commonly used when implementing locks
 - Acquire: no read/write after the load is reordered before
 - Release: no read/write before is reordered after this store
 - Relaxed: no restriction

Example of a simple mutex lock

```
pub struct Mutex { flag: AtomicBool, }
impl Mutex {
    pub fn new() -> Mutex {
        Mutex { flag: AtomicBool::new(false), }
    }
    pub fn lock(&self) {
        while self.flag.compare_exchange_weak(false, true,
            Ordering::Relaxed, Ordering::Relaxed).is_err()
            {}
        fence(Ordering::Acquire);
    }
    pub fn unlock(&self) {
        self.flag.store(false, Ordering::Release);
    }
}
```


d) Mutex

- Use `lock()` or `try_lock()` to access the data
 - Returns `Result<T>`
 - `lock()` is blocking mode
 - most usage simply `unwrap()` the result, why?
 - `try_lock()` is nonblocking mode
 - returns `Err()` if fails

```
let arc = Arc::new(Mutex::new(0));  
let arc1 = arc.clone();  
  
let _ = thread::spawn(move || -> () {  
    let mut data = arc.lock().unwrap();  
    *data += 1;  
}).join()
```

d) Mutex: Poison Strategy

- The lock could be poisoned if a thread holding the lock panics
- Use `into_inner()`

```
let arc = Arc::new(Mutex::new(0));
let arc1 = arc.clone();

let _ = thread::spawn(move || -> () {
    let mut data = arc1.lock().unwrap();
    panic!();
}).join()

assert_eq!(arc.is_poisoned(), true);
let mut guard = match arc.lock() {
    Ok(guard) => guard,
    Err(poisoned) => poisoned.into_inner(),
};
*guard += 1;
```

_____ The lock is poisoned

e) Condition Variable

- Synchronizing primitive used when threads need to wait for a resource to become available.

```
let pair = Arc::new((Mutex::new(false), Condvar::new()));
let pair2 = Arc::clone(&pair);
thread::spawn(move || {
    let (lock, cvar) = &*pair2;
    let mut started = lock.lock().unwrap();
    *started = true;
    cvar.notify_one();
});
let (lock, cvar) = &*pair;
let mut started = lock.lock().unwrap();
while !*started {
    started = cvar.wait(started).unwrap();
}
```

f) Once

- A synchronized primitive to run global initialization only one time
 - access `static mut` variables.

```
static mut VAL: usize = 0;
static INIT: Once = Once::new();

fn get_cached_val() -> usize {
    unsafe { INIT.call_once(|| {
        VAL = expensive_computation();
    });
    VAL
}
}
```

g) mpsc

- Multi-producer, single-consumer FIFO queue communication primitives
 - Asynchronous mode,
 - synchronous mode

```
let (tx, rx) = channel();
//let (tx, rx) = sync_channel()
for i in 0..10 {
    let tx = tx.clone();
    thread::spawn(move || { tx.send(i).unwrap(); });
}
//drop(tx); stop rx waiting
while let Ok(msg) = rx.recv()
    let j = rx.recv().unwrap();
    assert!(0 <= j && j < 10);
}
```


2. Marker Trait

Marker Traits

- Marker Traits have no methods to implement.
 - copy
 - sized
 - send
 - sync
 - unpin


Send

- The ownership of values of the type implementing Send can be transferred between threads.
- Use the move operator, which is similar as =
 - For types of Copy trait, make a copy of the object
 - For non-copy, transfer the ownership
- Almost all primitive types are Send
- Any struct composed of Send types is automatically marked as Send



Sync

- The type implementing Sync is safe to be referenced from multiple threads
- Any type T is Sync if &T is Send
- Sync is usually more rigid than Send
 - Are there any examples that are Sync but not Send?
 - cases are rare
 - exceptions may relate to thread-local features,
 - e.g., MutexGuard



Types Are Not Send

- Raw pointers are neither Send nor Sync
 - Possible to create shared objects (although unsafe)
 - Developers should tell compiler it is safe
- `RC<T>` cannot be Send, why?
 - Shared ownership
 - Problems in concurrent reference counter update
- `RefCell/Cell` are Send but not Sync
 - unsynchronized interior mutability

```
impl<T> !Send for Rc<T>  
impl<T> !Sync for Rc<T>
```

Can Cell be Send/Sync

```
fn check_sync<T: Sync>(_: T) {}
fn check_send<T: Send>(_: T) {}

fn testCell(){
    let mut v = Cell::new(1);
    check_send(v);           success
    //check_send(&mut v);    success
    //check_send(&v);        fail
    //check_sync(&v);        fail
    //check_sync(&mut v);    fail
}
```

Can Mutex be Send/Sync

- only if T is Send

```
fn testMutex1(){  
    let mut v = Mutex::new(1);  
    //check_send(v); success  
    //check_send(&v); success  
    //check_send(&mut v); success  
    //check_sync(&v); success  
    check_sync(&mut v); success  
}
```

```
fn testMutex2(){  
    let mut v = Mutex::new(Cell::new(1));  
    //check_send(v); success  
    //check_send(&v); success  
    //check_send(&mut v); success  
    //check_sync(&v); success  
    check_sync(&mut v); success  
}
```

Can Mutex be Send/Sync

```
fn testMutex3(){  
    let mut v = Mutex::new(&Cell::new(1));  
    //check_send(v);      _____ fail  
    //check_sync(v);      _____ fail  
}
```

```
fn testMutex4(){  
    let mut cell = Cell::new(1);  
    let mut v = Mutex::new(&mut cell);  
    check_send(v);        _____ success  
    //check_sync(&v);      _____ success  
}
```

Questions

- Does `ARC<T>` have bound on `T` to be thread-safe?
 - `T` does not need to be `Send + Sync` when constructing `ARC<T>`
 - The compiler checks the wrapped data during compilation

```
fn testArc(){  
    let mut cell = Cell::new(1);  
    let mut v = Arc::new(cell);  
    let v1 = v.clone();  
    thread::spawn(move || {  
        (*v1).set(3);  
    }).join();  
    (*v).set(2);  
}
```

Compilation error

Implementing Send/Sync is Unsafe

```
struct Unsend{ ptr: *mut i64, }
impl Unsend{
    fn add(&self, i:i64){
        unsafe{*(self.ptr) = *self.ptr + i};
    }
}
unsafe impl Send for Unsend{}
unsafe impl Sync for Unsend{}

fn main(){
    let mut var = 0i64;
    let mut v = Unsend{ptr:&mut var as *mut i64};
    let tid = thread::spawn(move || {
        for i in 1..100001{
            v.add(i);
        }
    });
    for i in 1..100001{
        var+=i;
    }
    tid.join();
    println!("{}",var);
}
```

In-Class Practice

- Rewrite your program (binary search tree or double-linked list) to be thread-safe
 - Support Sync/Send

More Reference

- <https://doc.rust-lang.org/nomicon/send-and-sync.html>
- <https://nyanpasu64.github.io/blog/an-unsafe-tour-of-rust-s-send-and-sync/>
- <https://cseweb.ucsd.edu/classes/sp17/cse120-a/applications/ln/lecture7.html>